

Parallelism in the commutation fold


From the algebra of series commutation to the execution schedule of a log-signature computation

signature-rs / lie-rs — 2026-09-02. Audience: readers comfortable with Lyndon bases, Lie series commutation, and log signatures; no familiarity with this codebase assumed. Concepts and methods only.

The three scales of the computation

A log signature of a path is built by folding the displacements of consecutive path points into a running Lie series, one Baker–Campbell–Hausdorff (BCH) concatenation at a time. The whole computation nests three scales, and the honest first answer to “where can threads go?” is different at each scale:

The path (scale 1)	An associative reduction of the per-segment displacements by BCH concatenation: $S = (((d_1 \oplus d_2) \oplus d_3) \oplus \dots)$. Computed today as a strictly sequential chain — fold $k + 1$ ’s input is fold k ’s output.	<i>serial between folds; associativity is the only escape hatch (§8)</i>
The fold (scale 2)	One concatenation $A \oplus B$ evaluates a fixed directed acyclic graph of Lie bracket nodes, compiled once from the BCH series truncated at the working degree. Nodes on one dependency level are independent; levels are dependency-ordered.	<i>parallel within each level; barriers between levels (§5)</i>
The commutation (scale 3)	Each node evaluates one series commutation $[X, Y]$ over the Lyndon basis: thousands of independent pair contributions, partitioned into anagram units that never share a target word.	<i>embarrassingly parallel inside a unit-partition (§2–4)</i>



Nesting: a serial chain of folds; each fold a level-ordered DAG of commutations; each commutation a bag of anagram units. Parallelism exists at the two inner scales; the outermost scale is a data dependency, not a scheduling choice.

The algebra: what one commutation computes

Brackets over the Lyndon basis

A Lie series $A = \sum_u a_u w_u$ is a formal sum over the Lyndon words — the canonical basis of the free Lie algebra on d letters. The commutator of two series distributes over pairs of basis words:

$$[A, B] = \sum_{u < v} (a_u b_v - a_v b_u) [u, v]$$

The sum runs over **canonical pairs** $u < v$ of Lyndon words (antisymmetry halves the work immediately). Each bracket $[u, v] = uv - vu$ re-expands in the Lyndon basis; the expansion has two rigid structural symmetries, and they are the entire foundation of the parallel layout:

Degree additivity	$ u + v = t$: every word in the expansion of $[u, v]$ has degree t . Truncating the series at degree m discards all pairs with $ u + v > m$ up front.
Content homogeneity	$[u, v]$ expands only over words with the same letter multiset as uv . Writing $\gamma(w)$ for a word’s content vector in $\mathbb{N}^d$: a pair with $\gamma(u) + \gamma(v) = \gamma$ can only ever touch basis words of content γ .

A **feasible decomposition** of a basis word w is a canonical pair (u, v) such that w appears (with nonzero integer coefficient) in the expansion of $[u, v]$, together with that coefficient $c_{u,v}^w \in \mathbb{Q}$. The coefficient of w in $[A, B]$ is then one scalar fan-in over all its feasible decompositions:

$$[A, B]_w = \sum_{(u,v) \rightarrow w} c_{u,v}^w \cdot \underbrace{(a_u b_v - a_v b_u)}_{\text{term — 1-2 multiplies}}$$

Each feasible pair contributes its bracket-coefficient times the antisymmetric product of the operand values. All $c_{u,v}^w$ are exact small rationals — precomputable once per basis, before any data arrives.



One word w ’s coefficient = a short inner product: its (pair, rational) decomposition row dotted with the antisymmetric pair-terms. Both the pair table and the rational rows are **structure constants**: they depend only on the basis, never on the data.

The precomputed structure-constant table

Because every ingredient except the a_u and b_v is data-independent, the entire commutation is driven by a single precomputed table:

Entries	All feasible canonical pairs (i, j) , packed as tiny fixed-size records, degree-blocked : sorted by (target degree t , then $p = u $, $q = t - p$). Degree blocking is what lets a truncated fold skip whole (p, q) regions of the table without testing them.
Decomposition rows	Per entry, the run of (target word, coefficient) pairs it feeds — the words whose expansion contains $[u, v]$. Most rows have length 1–3: almost every pair feeds its own distinguished word (the word uv itself when it is Lyndon) and almost nothing else.
Anagram-unit bounds	Markers (§3) delimiting the groups of entries that can write to the same set of target words — the unit is the conflict boundary for parallel writes.

Anagram classes and units: the partition that makes parallelism safe


Grouping by content

Group all degree- t basis words by their content vector γ . Each group — an **anagram class** — is the set of words obtainable from one another by permuting letters. Now consider a canonical pair (u, v) with $\gamma(u) + \gamma(v) = \gamma$, $|u| + |v| = t$: by content homogeneity, **everything it produces lands on the single class** (t, γ) .

The kernel groups entries by this pair — a (t, γ) -group is an **anagram unit** — and the payoff is the key structural fact of the whole engine:

Two distinct units never write to the same result word.

Each result word belongs to exactly one anagram class, and each unit feeds exactly one class. Therefore a whole commutation decomposes into units that are **disjoint write regions**: any two of them can be swept concurrently with no locks, no atomics, and no order constraints between them.



Measured shape of this partition (2 letters, degree ≤ 12): 747 basis words, 212 units, ≈ 1700 feasible pairs in one commutation; ≈ 3.5 words and ≈ 8 pairs per unit. At high degree the classes grow quickly (the class at content (6, 6) holds dozens of words), so units are numerous enough to fill many threads — but sized wildly differently, which matters for scheduling (§5).

Why a unit may never be split


Inside a unit, several canonical pairs feed the **same** result word (the word's decomposition row, §2). The result word accumulates its contributions sequentially — addition of floating-point values is not associative, so **the mathematical definition includes a fixed per-word accumulation order** (the order of the decomposition row). Two workers writing += into the same word would race (lost updates) and would break reproducibility even when they don't (interleaving changes the rounding). Hence:

The anagram unit is the atomic scheduling unit. Packs of parallel work are cut **only** at unit boundaries — a whole unit's entries always stay together, in table order, on one worker. Within a unit the sequential row order is honored by construction; across units no ordering exists at all. This is what makes results **bit-identical to the serial sweep at any thread count** — a property the codebase maintains as an invariant (and guards with cross-thread-count identity tests). The alternative — splitting units and compensating with atomics or privatized partial sums — costs either synchronization on the hot path or extra passes over memory, and an early version of the packer that accidentally cut inside units produced exactly the flaky wrong-results race this boundary exists to prevent.

One commutation, parallelized end to end

A single evaluation of $[X, Y]$ has four phases; two are parallel, two are serial-but-amortized. The architecture's recurring theme is: **anything whose value depends only on the operand supports (which words are nonzero), not on their values, can be precomputed and cached — supports change rarely, values change every fold.**

- | | | |
|------------------|---|---|
| ① Gating | From the two operands' supports, build presence bit-masks, then walk the entry table once to find active segments : the degree-runs whose (p, q) could produce anything. Output: per-unit active entry lists. | serial per call, cached by support fingerprint |
| ② Tickets | For each active segment, resolve once the per-entry presence tests (is u present in X ? v in Y ? both reversed?) and pack the survivors as compact tickets : the entry's table index plus two orientation bits. | part of the cached gating — not recomputed per fold |
| ③ Bundle | Cut the active units into a handful of packs of roughly equal total work, cutting only at unit boundaries. | cheap planning, parallel-safe by §3 |
| ④ Sweep | Workers claim packs dynamically. Per ticket: compute the antisymmetric term (1–2 multiplies) and scatter-multiply its decomposition row into the result. | the only heavy parallel region |



Ticket compaction in one segment: surviving entries (blue) become a dense ticket stream; entries that fail presence (grey, **X**) vanish from the sweep entirely. Measured at 2×12 : 31.7% of all entry visits were dead — tickets eliminate both the dead visits and the repeated bit-tests from every fold of the batch, for the price of one pass per support-change.

Two numbers characterize why these phases are arranged this way:

- **The sweep visits hundreds of thousands of entries per call** (measured $\approx 257,000$ tickets per fold at 2×12); the gating/ticket phases are $O(\text{table})$ once per support change. Across a batch of 1000 folds, the per-fold marginal cost of ① and ② is zero.
- **The serial fraction around a single call** (fork/join, task flattening, bundle build) is a fixed few microseconds — negligible against a millisecond sweep at 2×12 , but **larger than the entire kernel** at 12×2 (66 entries). This asymmetry returns in §7 as the work-adaptive scheduler.

The sweep itself

Per ticket the arithmetic is minimal: one or two multiplies for the term, then one fused multiply-add per decomposition-row element. What surrounds it is memory: the entry record, the operand words it indexes, the decomposed coefficients, and the scattered result words. Measured on real hardware, the per-entry time is **uniform** (11 ns/entry at 1 thread, ≈ 44 cycles) regardless of the pair's degree or row length — which is why balancing packs by entry count balances time (a fancier work metric was implemented, measured slower, and reverted) — and why the sweep parallelizes cleanly until **cache-coherence and streaming bandwidth**, not dependencies, set the limit (per-entry cost rises $\approx 13\%$ merely because two cores stream the same tables; $\approx 9 \times$ at 32 threads in the worst place — the accumulate traffic of §6).

The fold: a DAG of commutations with level barriers


From the BCH series to a DAG

Truncated at degree m , the BCH series $A \oplus B = A + B + \frac{1}{2}[A, B] + \frac{1}{12}[A, [A, B]] - \frac{1}{24}[B, [A, B]] + \dots$ has a fixed bracket structure — it depends only on m , not on data. The fold compiler turns this bracket forest into a DAG:

Nodes Bracket subterms, **interned** so shared subtrees (the same nested bracket can appear in many BCH terms) are computed once. Each node is one commutation kernel call — and, by content homogeneity, each node writes **exactly one anagram class** of target words.

Atoms The two operands: the running accumulator (A) and the folded-in displacement (B).

Levels A node’s level is the longest path from the atoms. Nodes of level ℓ read only results of levels $< \ell$, so each level is internally independent — many kernel calls with disjoint result buffers — while level $\ell + 1$ must wait for **all** of level ℓ .



The fold DAG is a **funnel**: thin at the atoms and at the deepest levels, widest in the middle (2×12: 11 levels, 600 nodes, peak 150 nodes at one level). Width per level is what a parallel schedule can harvest; the barriers between levels are what it must pay.


One parallel dispatch for the whole fold: the pack-slot walk

Naively, a fold with L levels pays L separate fork/join dispatches — measured at ≈35% of a small fold’s time. The folded engine instead plans **all** levels up front (legal because every node’s support is fixed for the fold’s duration) and dispatches **once**. Dependency order is enforced not by dispatch boundaries but by a small counter protocol:

Packs per level Each level’s (node, unit) tasks are cut into work-balanced packs — again, only at unit boundaries (§3). Because each node is one anagram class and tasks are node-ordered, every pack is **class-contiguous for free**, which keeps a pack’s scatter traffic inside a few cache-resident class blocks.

Slots A small number of long-lived worker tasks — never more than workers — each **claims packs off an atomic cursor** per level, executes them, then bumps the level’s completion counter.

Gates At each level boundary **every** slot waits for the previous level’s counter to reach its pack count (acquire/release ordering). A slot that arrives early is not allowed to sit idle **and**-blocking: before waiting, it **drains any unclaimed packs of the level itself**. That one rule is the whole liveness proof: a level’s remaining work is always either claimed-and-running or claimable by the waiter, so the walk cannot deadlock no matter how the pool is oversubscribed.



The walk across one fold (3 slots, 4 levels; shaded = sweep work, red = waiting). Barriers are implicit in the counters; empty levels or low-work phases are where slots park (and free their workers for co-scheduled work) after draining the cursor. Per-word accumulation order is exactly the serial schedule’s: packs never split a unit.

The batch: amortizing everything except the serial dependence

A path of 1000 segments is 1000 folds. Running them through the per-fold machinery would re-plan, re-gate, re-gather, and re-allocate 1000 times. The batched driver instead recognizes what is invariant **across** folds and pays it once:

Support fixed point After the first few folds the accumulator is dense; every node’s “which words does it write” list becomes a fixed point. From then on: no per-fold gating, no list rebuilds — the planner’s whole output (gates, tickets, packs) is reused verbatim for the remaining folds. Eligibility for batching is exactly this fixed point plus a shared displacement support.

Class-space accumulator The accumulator stays in the class-contiguous basis order across the whole batch (gathered once, written back once at the end), removing two full-basis permutations per fold.

Compact node buffers Each node’s result lives in a private scratch buffer holding only the degree slices it can touch (degree-blocked, shift-addressed), not the full basis — hundreds of nodes × full-basis buffers would be megabytes and cache-hostile; degree slices cut that 5×.

One stage chain, one walk The whole batch is a single chain of stages — per fold: **gather** the displacement (block-parallel), the L sweep levels, then **accumulate** the terms into the accumulator (block-parallel, simultaneously re-zeroing the node buffers for the next fold) — driven by exactly one pack-slot walk, sharing all cursors and counters.



The batch as one stage chain (5 of L sweep levels shown). Parallel work exists **inside** every stage; between stages, barrier counters; between the last accumulate of fold k and the first sweep of fold $k + 1$, a true data dependency — the accumulator. The chain is long and thin by design: the planning and allocation it replaces cost far more than the barrier hops when a fold’s sweep work is large, and far less (§7) when it is small.

The accumulate stage deserves one sentence of its own, because it is the batch’s second-largest cost: each fold’s terms must be added (with BCH weights) into the accumulator, and every node buffer re-zeroed, per fold. It runs block-parallel over class-position ranges — but measured element volume is large (≈575,000 slot-touches per fold at 2×12, ≈85% of the sweep’s own multiply count) because node buffers store whole degree slices while only support positions are live. Trimming the walk to live slots and fusing read-and-zero into one pass made this —63% smaller; what remains is **coherence-limited**: per-element cost grows 9× from 4 to 32 threads as the shared accumulator region crosses many cores — the one remaining place where more threads actively hurt inside a stage.

Scheduling to the shape of the work

The one structural number that decides whether parallelism helps is **planned sweep work per fold** — visible exactly, for free, inside the plan. The engine therefore sizes its own parallelism:

	planned entries per fold	chosen slots (32-thread pool)	e2e at 32t vs 1t
12×2 (high-dim, low-degree)	66	1 (serial walk)	2.55 → 1.25 ms ✓ (was 16.4 ms!)
8×3	700	1	28.6 → 7.7 ms ✓
3×8	≈74,000	19	338 → 161 ms ✓
2×12 (low-dim, high-degree)	≈257,000	32	1393 → 549 ms ✓
12×2 @32t	barrier/sync 92%		sweep
2×12 @32t	sync		sweep
			... acc

Same engine, two shapes, before work-adaptivity (illustrative shares, measured): at 12×2 the fold’s 66 entries vanish inside the stage machinery (16t = 6.6× **slower** than 1t); at 2×12 the sweep dominates and 16–32 threads pay. The scheduler’s rule — **slots ≈ planned work ÷ a fixed quantum, capped by the pool, floored at 1** — makes each shape run in the regime that suits it, at the cost of nothing when the work is large.


This single rule closed the high-dimension/low-degree pathology: the parallel infrastructure is still there, but a 66-entry fold no longer pays for it.

What threads can never reach — and the one axis they still could

It is worth stating plainly where parallelism **does not** exist, because it is a consequence of the mathematics rather than an implementation gap:

The fold chain is serial along the path. Fold $k + 1$ evaluates its DAG on the output of fold k ; level-0 jobs of fold $k + 1$ read the accumulator written by the accumulate stage of fold k . No reordering of stages, no matter how clever, removes a **data** dependency. Everything shown so far — units, levels, packs, slots, batches — lives strictly inside one concatenation, replicated identically 1000 times down the chain.

The one honest escape hatch already exists in the algebra: BCH concatenation is **associative**. The path’s log signature is an associative reduction of the segment displacements, so the reduction tree may have any shape — including a balanced tournament:



A tournament reduction over 8 segments: $\log_2 n$ sequential rounds, every pair in a round independent. Rounds start **sparse** (degree-1 displacements: cheap, support-gated folds) and only the last $\sim 2 \log_2 n$ concatenations along the critical path pay full dense-fold cost. A 1000-segment path goes from 1000 serial dense folds to 10 dense folds of dependency depth, with up to 500 independent concurrent concatenations in round 1.

Two caveats keep it out of the current engine: floating-point results would no longer be bit-identical to the sequential fold (reassociation changes rounding order — mathematically immaterial, contractually significant when tests assert exact bit equality), and the early rounds are so cheap that the same scheduling-discipline lessons of §7 apply at the tree level. But it is the **only** axis of parallelism that scales with path length — for the high-dimension/low-degree regime, whose per-fold work is intrinsically thin, it is the difference between “threads are politely declined” (today) and “threads map onto segments” (possible).

Regime map: what was measured

grid	stock, best config	today, best config	gain	why
2×12	1393 ms @32t	533 ms @16t	2.61×	All layers engaged: tickets, balanced packs, 16–19 slots of real work, adaptive accumulate. (<i>low-d, high-m</i>)
3×8	331 ms @8t	155 ms @16t	2.13×	Same, moderately narrower fold.
12×2	2.55 ms @1t (16t: 16.7 ms!)	1.25 ms — flat at 1t...32t	2.04×, pathology gone	66 entries/fold: adaptivity runs it serial-inside-parallel; batch amortization + allocation elimination. (<i>high-d, low-m</i>)
8×3	8.02 ms @2t (32t: 28.6 ms)	7.6 ms — flat	1.06×, pathology gone	700 entries/fold: right on the serial/parallel boundary.

End-to-end log signatures of 1000-segment random paths, criterion medians, 32-core machine. “Stock” = the engine before the parallelism work documented in the companion notes (unit-boundary pack cuts, one-dispatch fold, class-contiguous kernel, entry tickets, work-adaptive slots, allocation-free batch driver). Bit-identical results to the sequential fold throughout.

Summary. Parallelism in this engine is **structural**, not opportunistic: the algebra partitions work into anagram units that never share a write; the fold DAG partitions those units into dependency levels; the batch amortizes everything that is support-invariant; and the scheduler spends threads only where the measured work volume justifies the barriers. The serial spine of the computation — the fold chain along the path — is a data dependency of the BCH recursion itself, and the tree reduction of §8 is the only mechanism mathematics offers for crossing it.